

Programming #4 – Analysis Report

P1 – LLRB-based Multimap: My program `multimap.h` implements a multimap using a left-leaning-red-black tree data structure. Its public API includes `Size`, `Get`, `Contains`, `Max`, `Min`, `Insert`, `Remove`, and `Print`. The helper methods of `Get` (iterative), `Min`, `Insert`, `Remove`, `Print` (recursive), `IsRed`, `FlipColors`, `RotateRight`, `RotateLeft`, `FixUp`, `MoveRedRight`, `MoveRedLeft`, `DeleteMin` (self-balancing) aid in their implementation.

(1) I first changed the definition of a node by replacing a value with a deque of values and also rearranged the order, so the color & key come first in the initialization list. I decided to use a deque because the `Get()` & `Remove()` methods access the first element while `Insert()` pushes towards the end, thus, illustrating a queue form factor. I chose a deque over a queue because I iteratively `Print()` the multimap elements with `[]` access, which would also be more efficient than a list since nodes are contiguous in memory. (2) In `Get()`, I altered the line where the key would return the `front()` value of its list. (3) In `Insert()`, I first changed it so a new key at a `NIL` position would insert a new node with a key & deque of values with one value pushed back. I also changed it so when a key already exists, a new value will be pushed back. (4) In `Remove()`, I revised two locations where a key node is not at bottom & at bottom. For both cases I distinguished between `pop_front()` of the first element vs removing the entire node. When not at bottom, I made sure I replaced the node with the min node in the right subtree by swapping key-values. (5) In `Print()`, I printed out all key-value pairs of the internal deque.

In `test_multimap.cc`, I created several tests to increase the coverage of my implementation. Besides the given tests, `OneKey` & `MultipleKeys` which check an insertion of one key and multiple keys with `contains` & `gets`, `MaxMin` & `MultipleMaxMin` check the max & min of keys / duplicated keys with different values respectively. `ExceptionCase` checks that `get` should throw an exception when no key found. `DuplicateKeys` & `DuplicateKeyDiffValues` check duplicate insertion, removal, `get` of keys & keys with different values respectively; the latter also ensures proper print output. `GetDuplicateFirst` checks the first element obtained has the proper value with duplicated keys. `PrintInsertRemoval` & `RepInsertRemoval` insert & remove keys where the former ensures proper printing at the latter alternates between operations.

P2 – Completely Fair Scheduler: My program `cfs_sched.cc` receives a file of unordered task descriptions and feeds them into the CFS scheduler strategy. It uses the prior implementation of multimap.

The logic of `cfs_sched.cc` starts at the main method before flowing to respective methods which handle the various operations. The main method first performs error checking to ensure that a data file with tasks is passed in. It first calls `checkFileStream()` to make sure that the file is opened properly, and then `storeData()` to store the task data (id, start_time, duration) in a `Task*` vector. A `Task` class is made to hold its 3 primary variables along with runtime & vruntime and public accessors. Then, `organizeTasks()` is called to sort tasks with equal start times by id in alphabetical order. Finally, `runCFS()` is called which applies the CFS algorithm on tasks.

`runCFS()` is the primary function which calls the 7 steps of the CFS algorithm in partitioned public methods of the Scheduler class. The Scheduler holds a multimap of vruntime-`Task*` (key-value) pairs in its timeline. `appendTimeline()` first sees if any task start_time matches the tick value to be added to the multimap and sets its vruntime to the global value. `moveNextTask()` checks if the current task should move to the next task. `getNextTask()` gets the next task if current task is null and removes the next task from timeline. `incrementTask()` increments the current task's runtime & vruntime. `printStatus()` prints out the tick counter, total running tasks, and current task's id. `purgeCompletion()` deletes the current task if it is complete. `incrementTick()` increments the tick value for the next loop. Finally, `done()` is continuously checked to see if all tasks have been completed.